

The Heap and Destructors

Dynamic memory

Michael Burrell

May 2, 2024

Readings

Chapter 12 — 12.2 (*not* 12.1)

Objectives

- Learn how objects can be stored in two different regions of memory
- Solidify our understanding of the call stack
- Learn how to allocate objects on the heap
- Learn how to deallocate objects from the heap

Problems with the stack

- We have already seen some limitations with the call stack
- Notably, we cannot return a pointer to something in the current stack frame

```
1 int *foo() {  
2     int x = 5;  
3     return &x; // undefined behaviour!! bad news!!  
4 }
```

Motivations for the heap

- The call stack is managed for us automatically
 - New stack frame is created when we call a function
 - Current stack frame is destroyed when we return
- This is convenient, but limiting
- The heap is less convenient, but offers total freedom

Heap mechanism

- The heap mechanism is simple
 - Nothing happens implicitly
 - We must be explicit about all memory management in the heap
- Memory is created explicitly (using the `new` keyword)
- Memory is freed explicitly (using the `delete` keyword)

A note about Python

- Python uses the heap, as well, for objects
- However, Python has a feature called *automatic garbage collection*
 - In Python, memory does not need to be freed
- C++ does not have this feature, and memory management is done manually

Allocating a single object

- When allocating a single object (*not* an array), use the `new` keyword following by a constructor or type name
- In C++, even primitive data types can be constructed this way

```
1 int *returns_pointer_to_five() {  
2     int *r = new int;  
3     *r = 5;  
4     return r;  
5 }
```

Explanation of code

```
1  int *returns_pointer_to_five() {  
2      int *r = new int;  
3      *r = 5;  
4      return r;  
5  }
```

- The local variable `r` is allocated in the local stack frame
- However, it is pointing to a region of memory (big enough to hold an `int`) in the heap
- That memory in the heap will *not* be deallocated when this function returns
- We can return a pointer to this newly-allocated integer

Allocating an array

- Allocating an array on the heap is done similarly
- There is extra square bracket syntax needed to set the size of the array
- We get back a pointer to the first element of the array

```
1 int *x = new int[2];  
2 x[0] = 5;  
3 x[1] = 10;
```

Allocating arrays on the heap

- For large arrays or arrays where the size isn't known beforehand, we usually allocated on the heap
- On many operating systems, the call stack is of a limited size (often only a few megabytes)
 - The heap is big enough to store several billion gigabytes (limited by your computer's available RAM)
- For this reason, it's a common C++ idiom to allocate arrays on the heap unless they're small

Finally we can return an array!

```
1  int *nums_up_to(int max) {
2      int *r = new int[max];
3      for (int i = 0; i < max; i++) {
4          r[i] = i;
5      }
6      return r;
7  }
8  int main() {
9      int *array = nums_up_to(10);
10     cout << array[5] << endl; // will print out "5"
11     return 0;
12 }
```

Deallocating

- The code that we've seen so far in these slides has *memory leaks*
- A *memory leak* is when memory is allocated on the heap but never deallocated
 - It is especially bad if memory is allocated within a loop
 - The program's memory usage will keep increasing
- All memory allocated on the heap *should be* deallocated when it's no longer necessary

Example

```
1  int *nums_up_to(int max) {
2      int *r = new int[max];
3      for (int i = 0; i < max; i++) {
4          r[i] = i;
5      }
6      return r;
7  }
8  int main() {
9      int *array = nums_up_to(10);
10     cout << array[5] << endl; // will print out "5"
11     delete [] array; // this memory is not needed any
12                       more!
13     return 0;
14 }
```

Delete syntax

There are two forms of delete

`delete x` — when `x` points to memory that was allocated
as a single object

`delete [] x` — when `x` points to memory that was allocated
as an array

If you use the wrong one, your program will (probably) crash!

Dangling pointers

- The two most common mistakes when working with the heap are
 - memory leak** — when memory is not freed when it's not used any more
 - dangling pointer** — when memory *is* freed before it's finished being used

Dangling pointer example

```
1 int *x = new int;  
2 *x = 5;  
3 delete x;  
4 cout << *x << endl; // undefined behaviour!! bad news!!
```

Constructing and destroying objects

- Constructing and destroying objects is simple if you are comfortable already with allocating and deallocating primitive data

```
1  class student {  
2      string name;  
3  public:  
4      student(string name = "") : name(name) {}  
5  };  
6  int main() {  
7      student *s = new student("Joe Student");  
8      delete s;  
9      return 0;  
10 }
```

Default constructor

- Unlike in Python, the default constructor in C++ can be called without any parentheses

```
1 student *s = new student;  
2 delete s;
```

Destructors

- Things get more complicated when an object has a pointer to somewhere on the heap
- We have seen how to destroy the object itself, but not how to deallocate everything that the object is referring to

```
1 class student {  
2     double *grades; // what happens to the memory this is  
                       pointing to????  
3 public:  
4     student(int num_grades = 5) : grades(new  
        double[num_grades]) {}  
5 };
```

Destructors

- To get around this problem, C++ has the concept of *destructors*
- Destructors are *not* called explicitly
- They are called *implicitly* when an object *falls out of scope*
 - Stack-allocated objects will have their destructors called when their stack frame is destroyed
 - Heap-allocated objects will have their destructors called when the object is deleted

Example managing array

```
1  class student {
2      string name;
3      double *grades;
4      size_t num_grades;
5  public:
6      student(string name = "", size_t num_grades = 5)
7          : name(name), grades(new double[num_grades]),
8            num_grades(num_grades) {}
9      ~student() { // destructors start with a ~ tilde and
10                  cannot have parameters
11                  delete [] grades; // now we won't have a memory
12                      leak!
13      }
14  };
```

Rule of 3

- The *Rule of 3* is a convention followed in C++ object-oriented programming
 - Note that it is specific to C++ and not applicable to other OO languages
- It is used to manage memory in the heap
- It says “if any of the following 3 are implemented, then all 3 should be implemented”:
 - 1 destructor
 - 2 copy constructor
 - 3 assignment operator

Rule of 3 example

```
1  class pointer_to_int {
2      int *data;
3  public:
4      pointer_to_int() : data(new int) {} // default
5                                     constructor
6      ~pointer_to_int() { // destructor
7          delete data;
8      }
9      pointer_to_int(pointer_to_int const &o)
10         : data(new int) { // copy constructor
11             *data = *o.data;
12         }
13     pointer_to_int &operator=(pointer_to_int const &o) {
14         // assignment operator
15         *data = o.data;
16     }
```


Rule of 3

- As we see more complex data types, we will get more practice with the Rule of 3
- We have not used the copy constructor or assignment operator yet, but we will
- In all 3 methods, be sure to deallocate/allocate any memory necessary

Conclusion

- We can allocate memory on the stack or in the heap
- The heap is more free, and can store more memory
- The heap in C++ requires explicit allocation and explicit deallocation
- We can now create objects which are not bound to any particular stack frame