

I/O

Michael Burrell

## Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

# I/O

## C++ intro

Michael Burrell

March 5, 2024

# Readings for this set of slides

I/O

Michael Burrell

## Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

Chapter 1 — 1.2

Chapter 8 — 8.1

# Includes and namespaces

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- We're not ready to talk about `#include` and namespaces in detail yet
  - `#include` is coming up very soon
  - Namespaces will be explained even sooner
- We should at least understand them enough to get an idea of how “Hello world” works

# No standard prelude

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- Unlike many modern languages, C and C++ do not have a “standard prelude”
- When the compiler is run, it has *no* knowledge of *any* symbols at all
  - *Symbol* in this context refers to type definitions, classes, methods, objects, variables or constants
- This is in contrast to modern languages like Python
  - In Python, many symbols (e.g., `print`, `input`) are part of the “standard prelude”
  - They are implicitly imported into every Python file
- Note that *keywords* like `int` are not symbols and are known to the compiler innately

# Declaring foreign symbols

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- In C++, it is possible to declare foreign symbols (symbols defined in some other library, like the standard library) manually
- We will see how to do this in coming weeks
- It can get quite tedious to do this manually for every symbol we want to use
- The standard C++ library gives us *header files* (we can call them *include* files for now)
  - These files contain a list of some of the symbols that are defined in the standard library

# Two different types of includes

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- There are two different types of includes in C++
  - E.g., `#include "point.h"` with quotation marks
  - E.g., `#include <iostream>` with angle brackets
- Quotation-mark header files are header files that are defined within your project (not part of the standard library)
- Angle-bracket header files are header files that are external to your project (standard library or some 3rd party library)

# Our header file from the first week

## I/O

Michael Burrell

### Readings

Includes and  
namespaces

#include

Namespaces

### I/O

cout

cin

### Error handling

Life without exceptions

### Conclusion

`iostream` — part of standard C++. Defines all symbols needed for “I/O streams”

- We use I/O streams for console input/output, so this header file is very useful!
- It defines `cout`, `endl`, `cin`, and more

# Namespaces

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- Namespaces in C++ has some similarity to modules in Python
- They are a way to organize symbols (classes, variables)
- E.g., `cout` is a symbol which exists within the `std` namespace
- `std::cout` is the *fully-qualified name* of that symbol
- The `::` is used as a separator between namespace name and symbol name



# An example with fully qualified names in C++

I/O

Michael Burrell

Readings

Includes and namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

```
1 #include <iostream>
2
3 int main()
4 {
5     std::cout << "Hello_world" << std::endl;
6     return 0;
7 }
```

- Note the #include is necessary for bringing the cout and endl symbols into scope
- We still need to (by default) use the fully qualified name for any imported symbols

# Using using

I/O

Michael Burrell

Readings

Includes and namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

```
1 #include <iostream>
2 using namespace std;
3 int main()
4 {
5     cout << "Hello_world" << endl;
6     return 0;
7 }
```

- C++ offers a mechanism (called using) which brings in all symbols from a namespace
- It allows us to use the bare symbol name instead of the fully qualified name

# Using using

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

```
1 #include <iostream>
2 int main()
3 {
4     using namespace std;
5     cout << "Hello_world" << endl;
6     return 0;
7 }
```

- This is also legal (and common) in C++
- using follows the same scoping rules as variables
- In C++, we generally try to restrict using to the smallest scope reasonable

# include and namespace roundup

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

- Always use `#include <iostream>` and know that it's necessary for doing I/O
- Use or don't use `using namespace std;`
  - Know that it allows you to skip giving the fully qualified name of some symbols
  - Try to use it only when necessary, and in a small scope
  - Later on in the course, I will care more when and how you use it

# cout

## I/O

Michael Burrell

### Readings

Includes and namespaces

#include

Namespaces

### I/O

cout

cin

### Error handling

Life without exceptions

### Conclusion

- `cout`<sup>1</sup> is a standard object for printing character data to standard output
- It is an instance of the `ostream` class
- We have seen that we can use the `<<` operator to print things out with it
- Note that it makes no difference how we break up printing into separate operations

---

<sup>1</sup> “character out”

# Breaking up cout

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

---

```
cout << "hi" << 3 << endl  
    << 9 << "boo";
```

---

```
1 cout << "hi";  
2 cout << 3;  
3 cout << endl;  
4 cout << 9;  
5 cout << "boo";
```

---

These are equivalent.

# endl

## I/O

Michael Burrell

### Readings

Includes and namespaces

#include

Namespaces

### I/O

cout

cin

### Error handling

Life without exceptions

### Conclusion

- endl is a “basic I/O stream”
  - It is one of the types of objects which can be printed out with an ostream like cout
- It does two things:
  - 1 It prints out the newline ‘\n’ character
  - 2 It flushes the stream, just as if you did `cout.flush()`
- It is the preferred way to end lines in C++ because ‘\n’ by itself is not guaranteed to flush the stream

# Equivalence of methods and shift operator

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- By convention, we (almost) always use `<<` with `ostream` objects rather than calling methods directly
- If you find it clearer, you can usually use a method to do the same thing
- E.g., `cout << flush` is totally equivalent to `cout.flush()`
- We will generally follow common C++ convention of using `<<` everywhere



# Formatting

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- Formatting in C++ can be done with a few flags that are in the `std` namespace
  - `left`, `right`, `internal` — indicating horizontal alignment for upcoming item to print
  - `scientific`, `fixed` — which format to use when printing floating-point numbers
- We can also modify formatting with the following *methods*:
  - `width` — sets the width of the upcoming items to print
  - `precision` — number of digits after the decimal point, for floating-point numbers

# cin

## I/O

Michael Burrell

### Readings

Includes and  
namespaces

#include  
Namespaces

### I/O

cout  
cin

### Error handling

Life without exceptions

### Conclusion

- `cin`<sup>2</sup> is the counterpart to `stdout`
- It can read in items from standard input in a variety of different ways
- It is an object of the `istream` class

---

<sup>2</sup> “character in”

# cin

## I/O

Michael Burrell

### Readings

#### Includes and namespaces

#include

Namespaces

#### I/O

cout

cin

#### Error handling

Life without exceptions

#### Conclusion

---

```
1 int x, z;  
2 double y;  
3 cin >> x >> y >> z;
```

---

# ignore

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

- Much like `input` in Python, sometimes `cin`'s buffer has characters that you don't want to consider
- E.g., there might be a newline character hanging around that want to ignore
- `cin` has an `ignore` method that will allow you to do this

# ignore

## I/O

Michael Burrell

### Readings

Includes and  
namespaces

#include  
Namespaces

### I/O

cout  
cin

### Error handling

Life without exceptions

### Conclusion

---

```
1 int x;  
2 string z;  
3 cin >> x;  
4 cin.ignore(SSIZE_MAX, '\n');  
5 cin >> z;
```

---

ignore has two parameters: the maximum number of characters to ignore, and the type of character to stop ignoring at.

# numeric\_limits

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include

Namespaces

I/O

cout

cin

Error handling

Life without exceptions

Conclusion

- On the previous slide, I used `SSIZE_MAX` as the maximum number of characters to ignore
- We want to say “ignore *any* number of characters up to the next newline”, which has to be stated as “ignore the maximum number of characters up to the next newline”
- `SSIZE_MAX` will *probably* be correct, but in production code, you should use `numeric_limits<streamsize>::max()` instead
  - This is using an advanced C++ feature called templating
  - In this course, you are permitted to use `SSIZE_MAX` instead

# Checking return values

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

- C++ *has* exceptions
- C++ programmers almost never use them
- Exceptions have a performance cost associated with them
- Instead of relying on exceptions, we check return values to see if something worked

# Did cin work?

## I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

---

```
1  int x;
2  cin >> x;
3  if (cin) {
4      cout << "You entered in" << x;
5  } else {
6      cout << "That was not a valid integer";
7      cin.clear();
8      cin.ignore(SSIZE_MAX, '\n');
9  }
```

---

We can treat cin as if it were a bool to check if it's in an *error state* or not.



# Error states

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

- If an `istream` enters an error state (due to bad input, end of file, keyboard unplugged, etc.) it *will not work* until its error state is *cleared*
- We can clear an error state using the `clear` method

# Clearing the buffer

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

- The `clear` method does *not* clear the buffer
- The `clear` method does NOT clear the buffer

■ **The `clear` method  
does NOT clear the  
buffer**

- The `clear` method only transitions `cin` from an error state to a normal state
- Use the `ignore` method after the `clear` method to clear out the input buffer

# Repeatedly reading in

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

## Class exercise

Let's write a C++ program which repeatedly asks the user for input until the user enters in a number.

# What we learned

I/O

Michael Burrell

Readings

Includes and  
namespaces

#include  
Namespaces

I/O

cout  
cin

Error handling

Life without exceptions

Conclusion

- Includes are necessary for bringing in symbols
- using namespace is sometimes convenient for not giving a fully qualified name of a symbol
- Formatted output in cout
- Input with cin
- Use if (cin) to check if input was read successfully, then clear+ignore if it wasn't